

SUDHARSAN R

AI Product Manager

Chennai, India ·

sudharsan454@gmail.com · [linkedin.com/in/sudharsan454](https://www.linkedin.com/in/sudharsan454)

SUMMARY

AI Product Manager with 9 years in B2B enterprise SaaS. I own two production AI products at Facilio, Helpdesk AI and Payment Validator, reporting to the CEO: strategy, roadmap, evaluation, pricing, and revenue. My edge is problem selection: I sit inside customers' operations, find the one workflow where AI genuinely pays for itself, and prove it in a controlled pilot before expanding. I build the prompts, eval harnesses, and test tooling myself, so the loop from customer conversation to shipped agent change runs in days, not quarters.

CAREER HIGHLIGHTS

- ~80% of helpdesk requests resolved end-to-end by AI at a flagship enterprise deployment: 24/7 first responder, live across the UK, Middle East, and ANZ.
- Took the AI line from launch to ~\$200K ARR across 5 paying customers in its first 12 months, with 2 paid and 4 free pilots currently in flight.
- Shipped the Field Service Management vertical 0→1 in 5 months with a team of 15, unlocking \$1M TCV potential in year one.
- \$4.5M+ ARR contributed across CMMS, FSM, and retail; rescued a ~\$400K POC on-site with the customer's US leadership.

EXPERIENCE

Facilio, Chennai — *Enterprise SaaS for real-estate & facilities operations (CMMS · FSM · AI)* | **Oct 2021 – Present**

AI Product Manager | Nov 2025 – Present

- Own Helpdesk AI and Payment Validator end to end, reporting to the CEO across strategy, roadmap, evaluation, pricing, and GTM; grew the AI line to ~\$200K ARR across 5 paying customers in its first 12 months, with 2 paid and 4 free pilots in flight across the UK, Middle East, and ANZ.
- Flagship deployment runs an entire enterprise helpdesk line 24/7 with AI as the first responder: ~80% of requests diagnosed, routed, and resolved end to end; the remaining 20% escalated to humans through a handoff path designed in from day one.
- Own pricing end to end: trialled and killed per-seat and advance-credit models before landing per-action usage pricing (per voice-minute, per conversation, per validated invoice) plus outcome-based pricing per work order dispatched where the outcome is measurable; brought existing customers along through each shift.
- Discovery at scale: 80 structured discovery calls in 4 months surfaced the two universal choke points — service intake and invoice approvals — and that insight set the AI roadmap: Helpdesk AI first, Payment Validator next.
- Build hands-on: prompts, eval harnesses, and test tooling written myself with Claude, Cursor, and MCP; the loop from customer conversation to shipped agent change runs in days, not quarters. Agents ship on single-agent, agent-handoff, and orchestrator patterns with RAG over customer knowledge bases.
- Made evaluation the release gate: 50+ real scenarios per agent, every round scored against actual backend records, with emergencies, mixed intents, junk input, and data-extraction attempts tested before any customer touched the agent; agents ship with known failure modes, not surprises.
- Make the trade-off calls that decide whether an AI product survives real operations: response speed vs. reasoning depth per channel (a tenant on a voice call won't wait; an email agent can think longer), model and token cost per workflow, and human-in-the-loop boundaries where customers drew them, aligning engineering on model choice, fallback behavior, and which guardrails live in the prompt vs. the pipeline.
- Convert enterprise interest into signed pilots through tightly scoped SOWs (one use case, one site, hard success criteria), backed by CXO demos and competitive positioning against incumbent vendors; run the demand-gen motion alongside the CPO through webinar-sourced enterprise leads, BDR plays, and a behavior-triggered self-serve trial funnel.

Lead Product Manager | Jan 2025 – Oct 2025

- Led 0→1 launch of the Field Service Management (FSM) vertical for retail chains and FM companies: shipped in 5 months with a cross-functional team of 15, unlocking \$1M TCV potential in year one.
- Rescued a ~\$400K POC through fast iteration, direct stakeholder engagement, and on-site whiteboarding with the customer's leadership in the US.

- Owned the FSM and retail verticals end to end within a \$4.5M+ ARR portfolio; set FSM pricing with distinct denominators per segment; mentored 4 APMs owning future product bets.

Product Manager | Oct 2021 – Dec 2024

- Owned the core CMMS product (~\$4.5M ARR, 25+ enterprise customers); cut bugs 30% in a critical module by grounding fixes in direct customer sessions instead of feature requests.
- Grew multi-module adoption ~10% in 12 months → ~\$50K upsell and prevented churn on two accounts worth \$330K ARR; contributed to 8% YoY ARR growth from expansion plays.
- Main product PoC for pre-sales: 20+ SOWs reviewed, 15 BRDs, 40+ RFP/RFI submissions in 2 years.
- Shipped a self-service desktop app for bulk data operations, adopted by all customers, cutting onboarding friction.

Tata Consultancy Services, Chennai — *Product Manager* | Aug 2019 – Sep 2021

- Built a blockchain SaaS platform for digital bank guarantees (TCS Quartz) serving 15+ issuers and 5+ government entities; owned requirements, workflow design, and launch scope across bank and government stakeholders, cutting issuer workload by 5 FTEs; first deployment signed by Israel's largest bank.

Accenture, Chennai — *Test Engineering Analyst* | Jun 2015 – Aug 2017

- QA for BofA's Liquidity Risk Management application (BASEL III, 100+ upstream sources); regression suite overhaul lifted defect identification 15%.

SKILLS

- **AI product judgment:** problem selection from live operations, latency/accuracy/cost trade-offs, human-vs-AI boundary design, evaluation frameworks as release gates, pilot scoping that converts
- **Product core:** product strategy, product discovery & customer interviews, roadmap prioritization, PRDs/BRDs/SOWs, stakeholder management, Agile delivery with cross-functional teams, success metrics & KPIs, product analytics (PostHog)
- **AI/LLM:** LLM agent design, generative AI, agentic AI, RAG, multi-agent orchestration, conversational & voice AI, prompt engineering, LLM evaluation (Langfuse), MCP, human-in-the-loop design
- **Commercial:** usage-based & outcome-based pricing and packaging, enterprise GTM, competitive positioning, CXO demos
- **Working stack:** Claude, Cursor, MCP, Langfuse, Facilio AI Studio

Education: MBA, University of Madras (2017–2019) · B.E. Mechanical Engineering (2015)